

Primary analysis results, n=13 (run once, 2 September 2026)

I ran the preregistered pipeline ([src/analyze.py](#)) once on the 13 real CSVs. All of them came in on task-v2.3-FROZEN, and everyone passed both catch trials, so nobody was excluded. The raw output is archived at [data/analysis_output_2026-09-02.txt](#).

The verdict

EV beat Adaptive Gain by more than the preregistered margin of 3 log units, shown in the table. The sensitivity analysis agreed with the primary result on both ranking and margin. The EV/Adaptive Gain pair beat the other models by a wide margin, which makes sense because they are mathematical cousins.

model	w	beta	LL prim	CV-LL	LL full
EV	0.61	18.9	-349.3	-352.2	-440.9
AdaptGain	0.59	10.5	-353.6	-355.5	-444.2
Divisive	0.50	6.0	-358.3	-360.3	-444.7
Recurrent	0.50	21.5	-360.4	-363.9	-447.4
PT	0.59	5.4	-378.2	-379.9	-464.5
Range	0.51	3.4	-390.4	-391.5	-481.9

What people actually did

The baseline split came out at 47% A (49 of 104), which means the pilot titration pretty much hit the 50/50 target. But the cohort was pretty polarised. Six people chose B almost all the time and five chose A pretty much all the time, and two others were more flexible. So the balance comes from committed A-people and committed B-people cancelling out rather than from anyone sitting on the fence. The model also shows that people's choices were very consistent.

The dominance check went how you'd expect. The pattern followed preference almost perfectly, in that the A-people almost always took the upgraded version of A, while the B-people went right past a deal that didn't dominate anything in their eyes.

cell	n(A/B)	A share	95% CI	decoy-picked
d1	22	5%	[1, 22]%	66%
d2	27	19%	[8, 37]%	58%
d3	65	40%	[29, 52]%	0%
d4	65	43%	[32, 55]%	0%
d5	62	45%	[33, 57]%	5%
d6	64	44%	[32, 56]%	2%

Willingness-to-pay slopes mostly fell between £2.80 and £7.50 per doubling of data, with IS (13.0), MA (9.6) and al (20.8) running high.

Why the table's anomaly isn't a context effect

The number that is most likely to be misinterpreted is the 5% conditional A-share at d1, next to a 47% baseline. It looks like the decoy shoved everyone toward B, but it didn't. What happened is selection, because the A-preferrers took the dominating decoy and in that left the conditional pool, so the A-vs-B trials that remain at d1 come almost entirely from the B-preferrers, who were never going to pick A anyway. The true test for decoy influence is at the four non-dominating cells, where nobody is being filtered, and there the conditional shares (40 to 45%) sit right on top of baseline. So the model-free picture and the fitted verdict are telling one story between them. Thirteen decisive people with stable, but very different price/data preferences chose whatever their own values said to choose, and the decoys moved nobody we can detect.

Limitations and Considerations

Thirteen people instead of the planned eighteen (I had a few non-respondents). The fitting is pooled at the cohort level, like in the preregistration, even though baseline behaviour is strongly bimodal across individuals. And EV's margin over its cousin Adaptive Gain is right at the cutoff. Anything I do next that splits people by preference group is exploratory and will be labelled as such.

Next

I will do a results write-up, winning-map figure, possibly an exploratory per-participant split, and a results one-pager to present.